

STA 5635 HW 7

Palmer Swanson

February 22, 2026

Problem 1

(1a) NN w/ one hidden layer, 256 neurons, ReLU activation.




Figure 1: Train and Test accuracy vs. Epoch, part a.

(1b) NN w/ two hidden layers, 64 neurons, ReLU activation

(1c) CNN w/ one hidden layer, 32 filters of size 15, ReLU activation

(1d) CNN w/ one hidden layer, 32 filters of size 15, ReLU activation, max pooling size 2, stride 2.

(1e)